

Fundamentals of deep learning

Crash course in using PyTorch to train models



Terence Parr

MSDS program

University of San Francisco

The goal of (supervised) machine learning

- Obtain a model that predicts numeric values (regressor) or categories (classifier) from information about an entity:



The abstraction

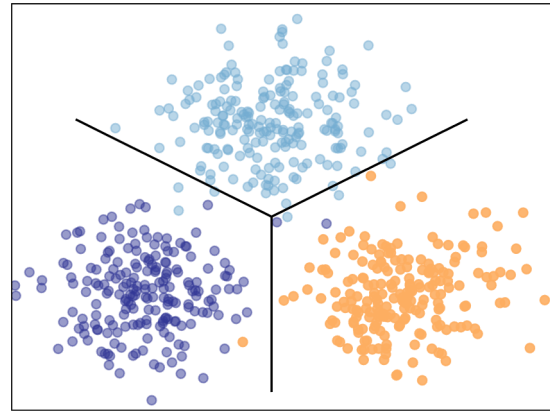
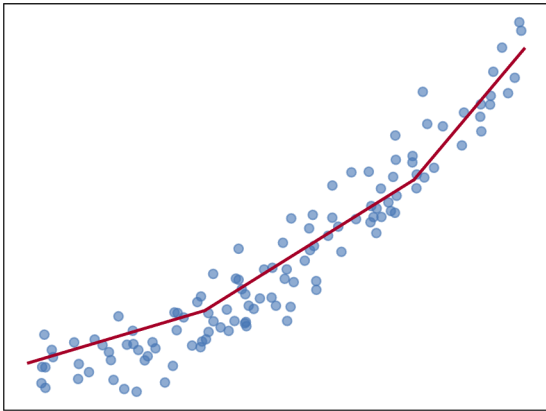
- Training a model means capturing the relationship between feature vectors and a target variable, given a training data set
- **Feature vector** x : set of features or attributes characterizing an entity, such as square footage, num of bedrooms, bathrooms
- **Target** y : either a scalar value like rent price (regressor), or an integer indicating “creditworthy” or “it's not cancer” (classifier)
- Model captures this relationship: $X \rightarrow y$

X	y
\mathbf{x}_1	y_1
\mathbf{x}_2	y_2
\vdots	\vdots
\mathbf{x}_n	y_n

Row vectors, x_i ,
represent instances

Regression versus classification

- Predictors fit curves to data and classifiers draw decision boundaries between data points in different categories
- Two sides of the same coin in implementation typically



Models

- Models are composed of *parameters*; predictions are a computation based upon these parameters
- Models have *architecture*; e.g., number of layers, number of neurons per layer, which nonlinearity to use, etc...
- Models have *hyper-parameters* that govern architecture and the training process; e.g., learning rate, weight decay, drop out rate
- Hyper-parameters are specified by the programmer, not computed from the training data; must be tuned
- Deep learning training greatly affected by learning rate, and even things like parameter initialization

Training

- Training a model means finding optimal (or good enough) model parameters as measured by a *loss* (cost) function
- Loss function measures the difference between model predictions and known targets
- *Underfitting*: model unable to capture the relationship $X \rightarrow y$ (assuming there is a relationship to be had)
- *Overfitting*: model is too specific to the training data and doesn't generalize well (e.g., home price predictor trained on just condos)
- To *generalize* means we get accurate predictions for test feature vectors not found in the training set

Terminology: Loss function vs metric

- *Loss function*: these are minimized to train a model
E.g., gradient descent uses loss to train regularized linear model
- *Metric*: evaluate accuracy of predictions compared to known results (the business perspective)
- Both are functions of y and \hat{y} , but loss is also possibly model parameters (e.g., linear model regularization loss tests parameters)
- Examples:
 - Train: MSE loss & Metric: MSE metric
 - Train: MSE loss & Metric : MAE metric
 - Train: Negative Log Loss & Metric : misclassification or FP/FN metric
- If metric is applied to validation or test set, informs on generality and quality of your model

See also stackoverflow post by Chstiros Tsatsoulis: <https://goo.gl/T5AmrT>

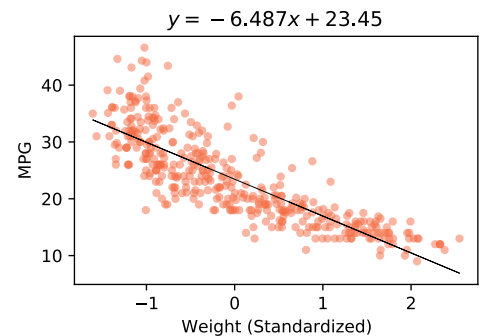
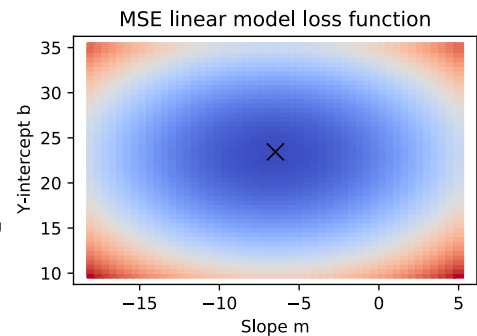
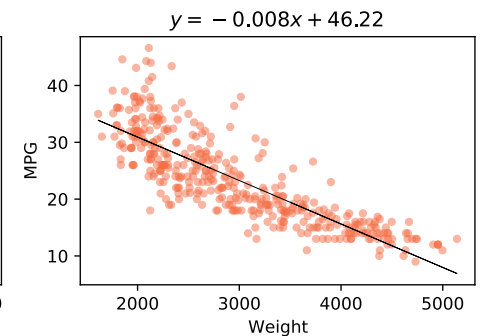
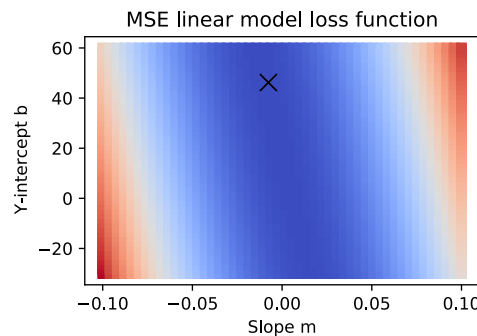
Train, validate, test

- We always need 3 data sets with known answers:
 - training
 - validation (as shorthand you'll hear me & others call this test set)
 - testing (put in a vault and **don't peek!!**)
- Validation set: used to evaluate, tune models and features
 - Any changes you make to model tailor it to this specific validation set
- Test set: used exactly **once** after you think you have best model
 - The only true measure of model's generality, how it'll perform in production
 - Never use test set to tune model
- Production: recombine all sets back into a big training set again, retrain model but don't change it according to test set metrics

Preparing data

X	y
\mathbf{x}_1	y_1
\mathbf{x}_2	y_2
\vdots	\vdots
\mathbf{x}_n	y_n

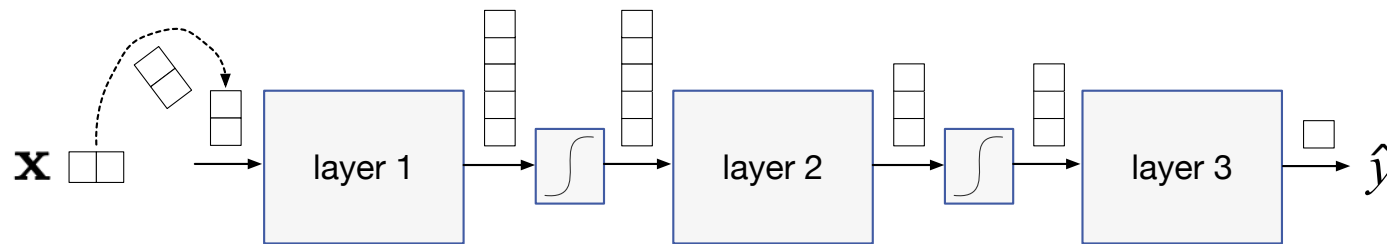
- Everything must be numeric
- No missing values
- Dummy encode categoricals
- Should normalize numeric features in X to zero-mean, variance one ("whitening")
- Speeds up training
- Compare regression equations, loss function surfaces



Deep learning regressors

What's a neural network?

- A combination of linear and nonlinear transformations
 - Linear: $z^{[layer]} = W^{[layer]}x^T + b^{[layer]}$
 - Nonlinear example: $a^{[layer]} = \sigma(z^{[layer]})$; called *activation function*
- Ignore the neural network metaphor, but know the terminology
- Networks have multiple *layers*; layer is a stack of *neurons*



- Transform raw x vector into better and better features, final linear layer can then make excellent prediction

DL Building blocks

$$\mathbf{w}\mathbf{x} + b$$

linear





sigmoid



ReLU (rectified linear unit)

- $\hat{y} = w_1x_1 + w_2x_2 + \dots + w_mx_m + b = \mathbf{w}\mathbf{x}^T + b$ for $n \times m$ dim X
- Linear/logistic regression equivalents (one x instance)



Underfitting a bit here
(need more of a quadratic)

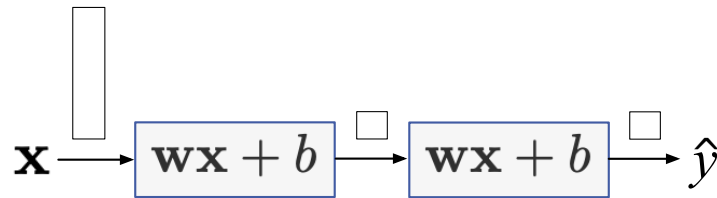
Assume we magically know \mathbf{w} and b

Try adding layers to get more power

- But, sequence of linear models is just a linear model

$$\hat{y} = w'(\mathbf{w}\mathbf{x}^T + b) + b' = w'\mathbf{w}\mathbf{x}^T + w'b + b' = \mathbf{w}''\mathbf{x}^T + b''$$

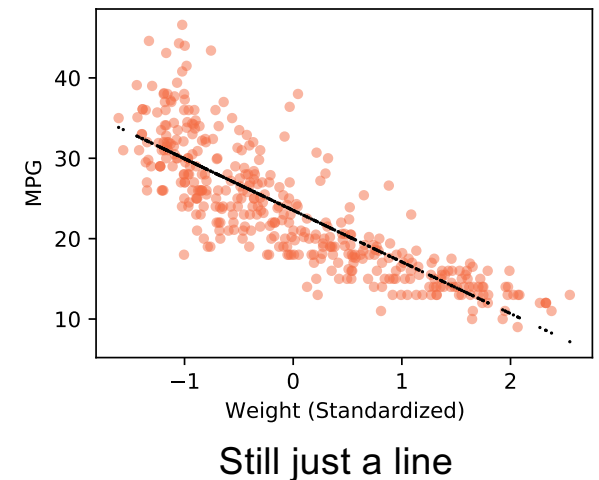
(w' is scalar since $\mathbf{w}\mathbf{x}^T + b$ is scalar)



- PyTorch code

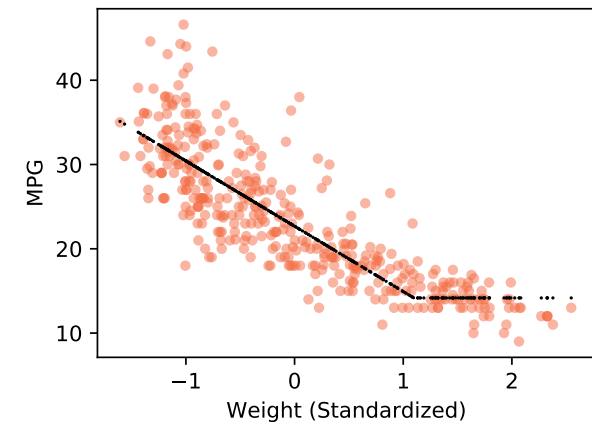
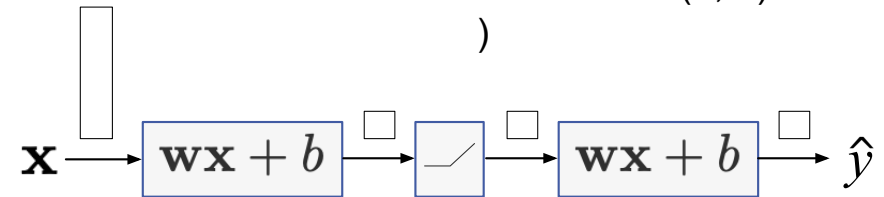
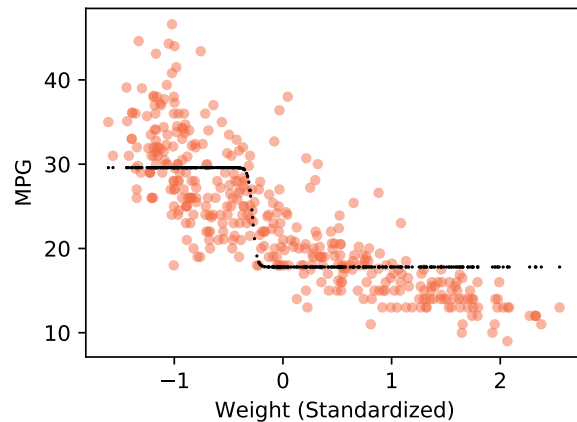
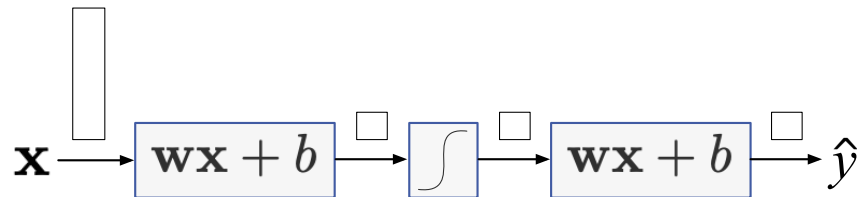
```
model = nn.Sequential(  
    nn.Linear(m, 1), # m features  
    nn.Linear(1, 1)  
)
```

(For simplicity, I'm using proper $\mathbf{w}\mathbf{x}^T$ in math but omitting transpose in diagrams)



Must introduce nonlinearity

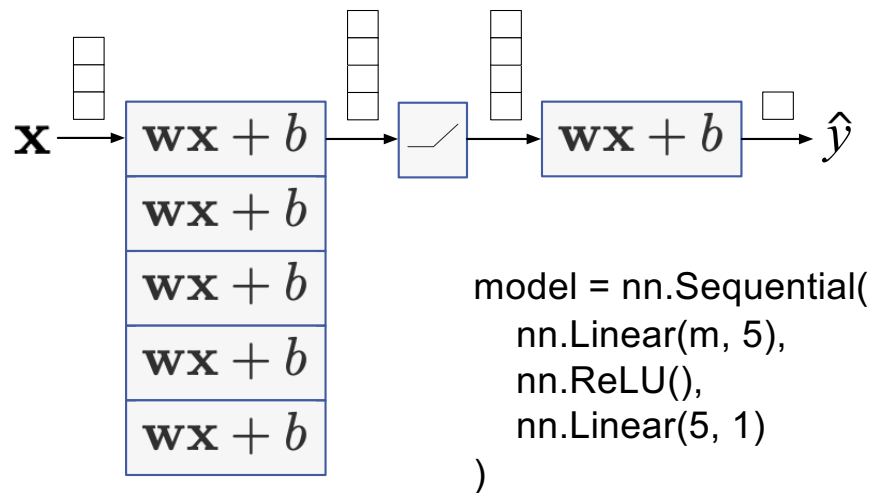
```
model = nn.Sequential(  
    nn.Linear(m, 1),  
    nn.ReLU(),  
    nn.Linear(1, 1)  
)
```




ReLU idea here: Draw two lines
then clip at intersection

Stack linear models (neurons) for more power

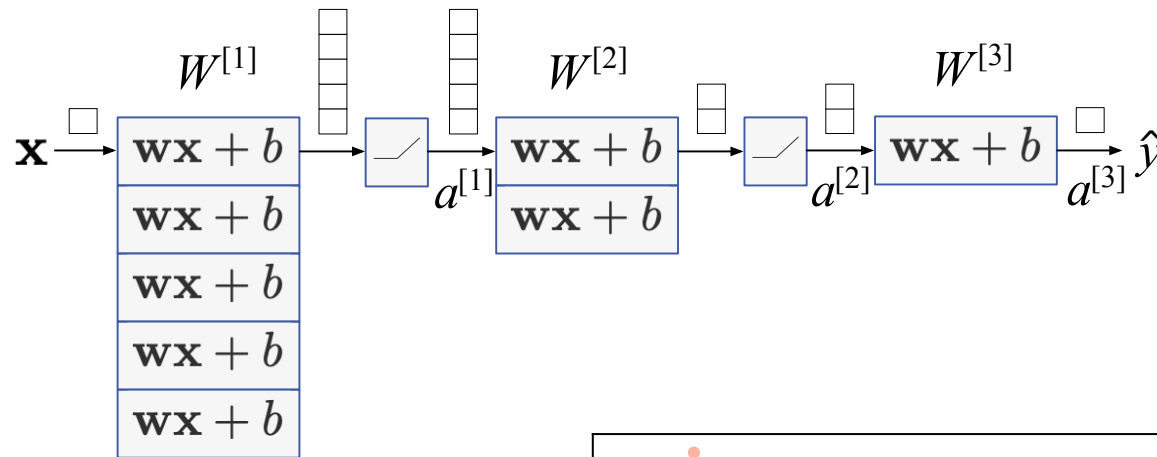
- Stack gives layer: W matrix and b
- $a^{[1]} = \text{relu}(W^{[1]}x^T + b^{[1]})$
- $\hat{y} = a^{[2]} = W^{[2]}a^{[1]} + b^{[2]}$



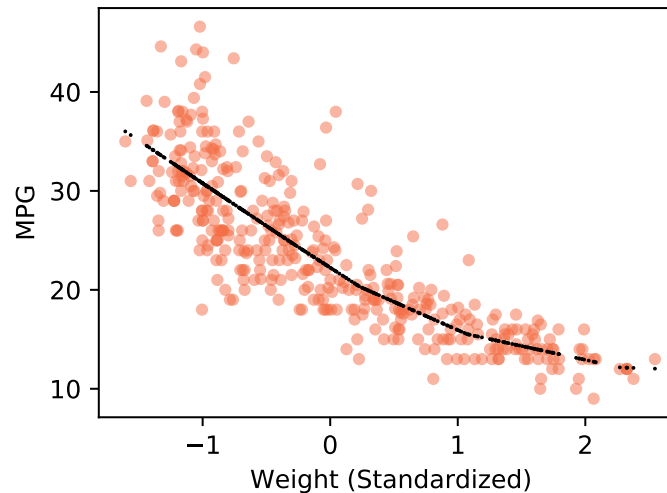
All those w and b are different



Math for Cars dataset 1D: weight→MPG



```
model = nn.Sequential(
    nn.Linear(1, 5),
    nn.ReLU(),
    nn.Linear(5, 2),
    nn.ReLU(),
    nn.Linear(2, 1)
)
```



$$a1 = F.relu(W1 @ x)$$

$$a2 = F.relu(W2 @ a1)$$


$$a3 = W3 @ a2$$

(courtesy of TensorSensor)

Too much strength can lead to overfitting

- Models with too many parameters will overfit easily, if we train a long time
- We'll look at regularization later

```
model = nn.Sequential(  
    nn.Linear(1, 1000),  
    nn.ReLU(),  
    nn.Linear(1000, 1)  
)
```

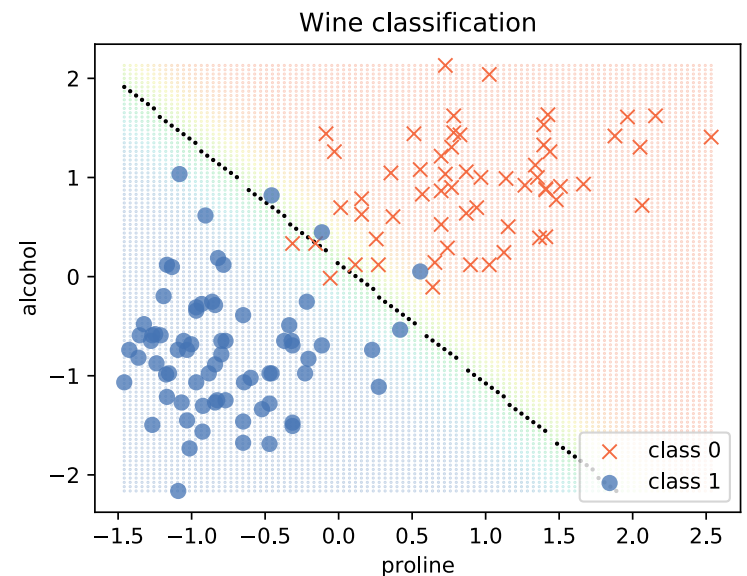
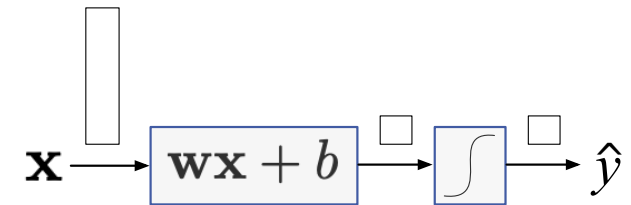


Classifiers

Binary classifiers

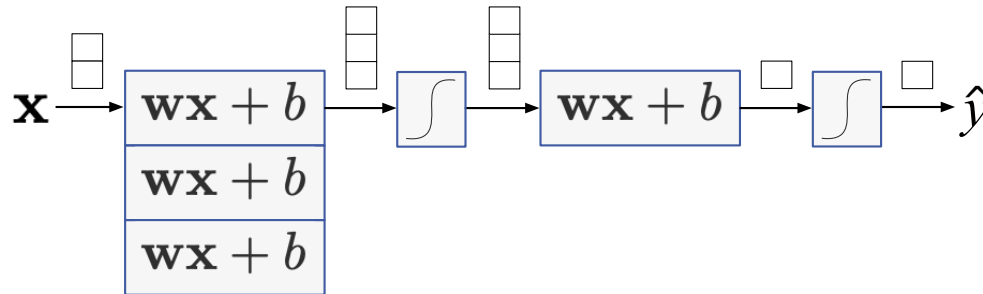
- Add sigmoid to regressor and we get a two-class classifier
- Prediction \hat{y} is probability of class 1
- One-layer network with sigmoid activation function is just a logistic regression model
- Provides hyper-plane decision surfaces

```
# 2 input vars: proline, alcohol  
model = nn.Sequential(  
    nn.Linear(2, 1),  
    nn.Sigmoid(),  
)
```




Stack neurons, add layer

- We get a nonlinear decision surface

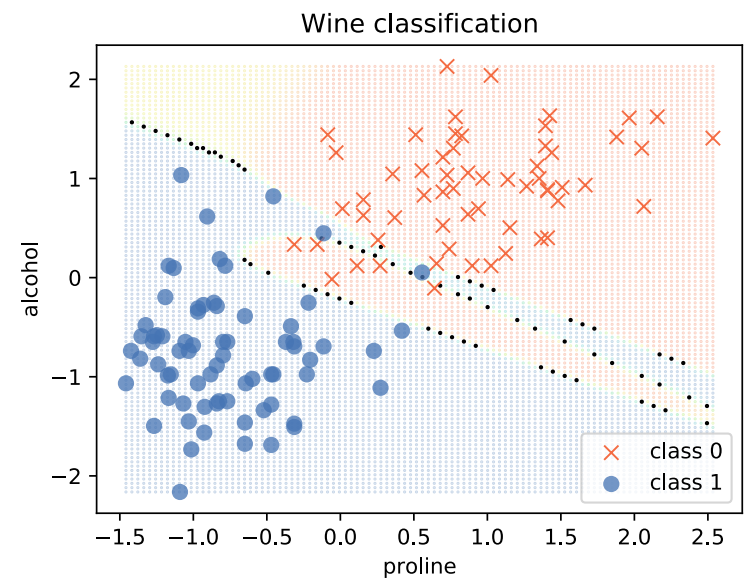
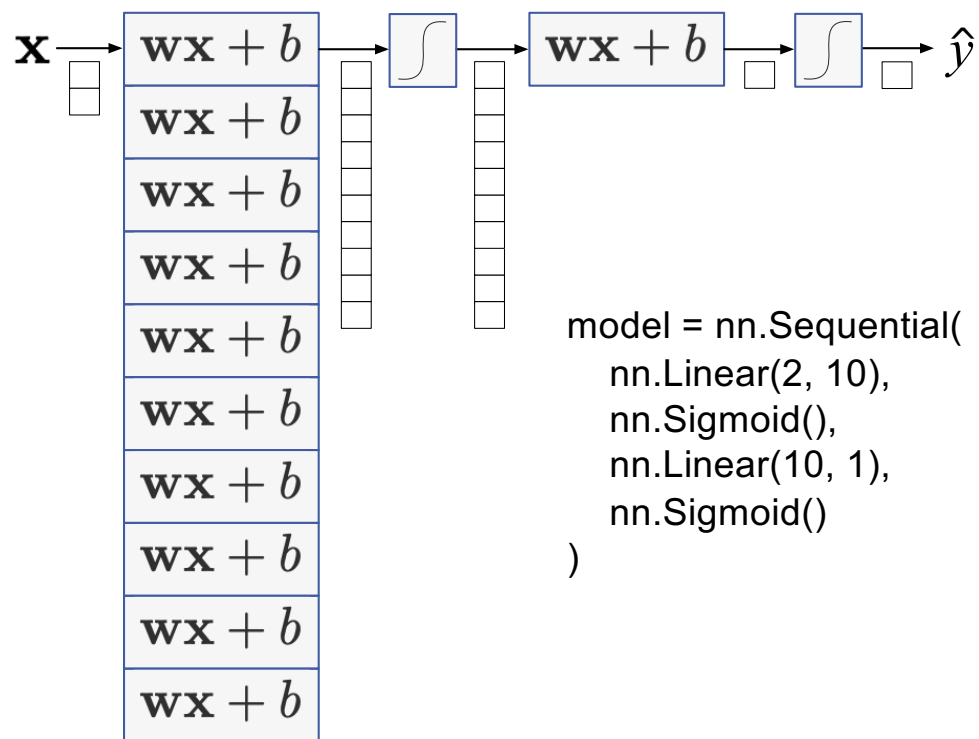


```
model = nn.Sequential(  
    nn.Linear(2, 3),  
    nn.Sigmoid(),  
    nn.Linear(3, 1),  
    nn.Sigmoid()  
)
```

All those w and b are different



More neurons: more complex decision surface

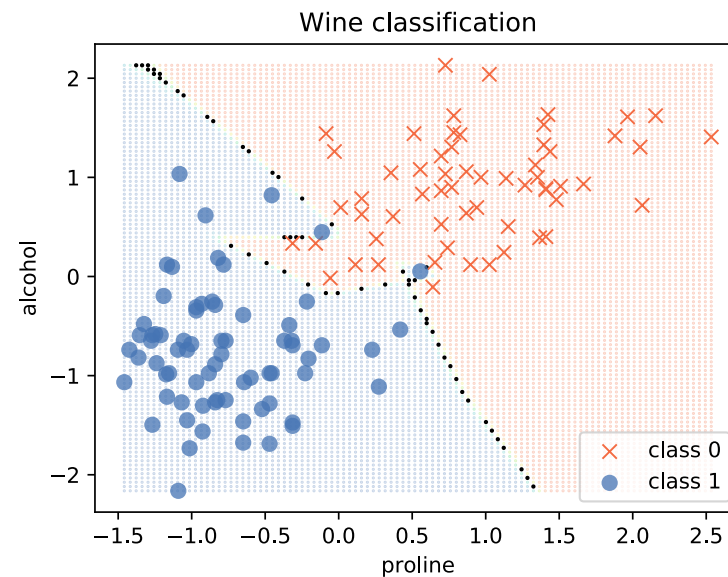


Likely overfit

Even ReLUs can get curvilinear surfaces

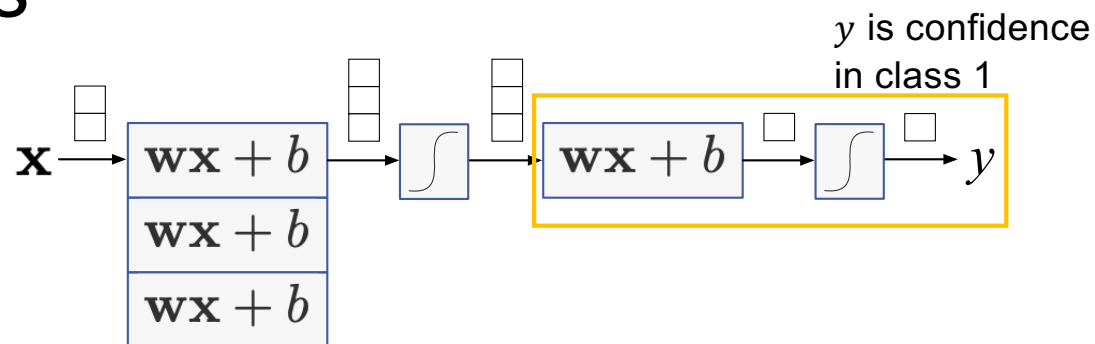
```
model = nn.Sequential(  
    nn.Linear(2, 10),  
    nn.ReLU(),  
    nn.Linear(10, 10),  
    nn.ReLU(),  
    nn.Linear(10, 1),  
    nn.Sigmoid()  
)
```

(Last activation function still must be sigmoid)

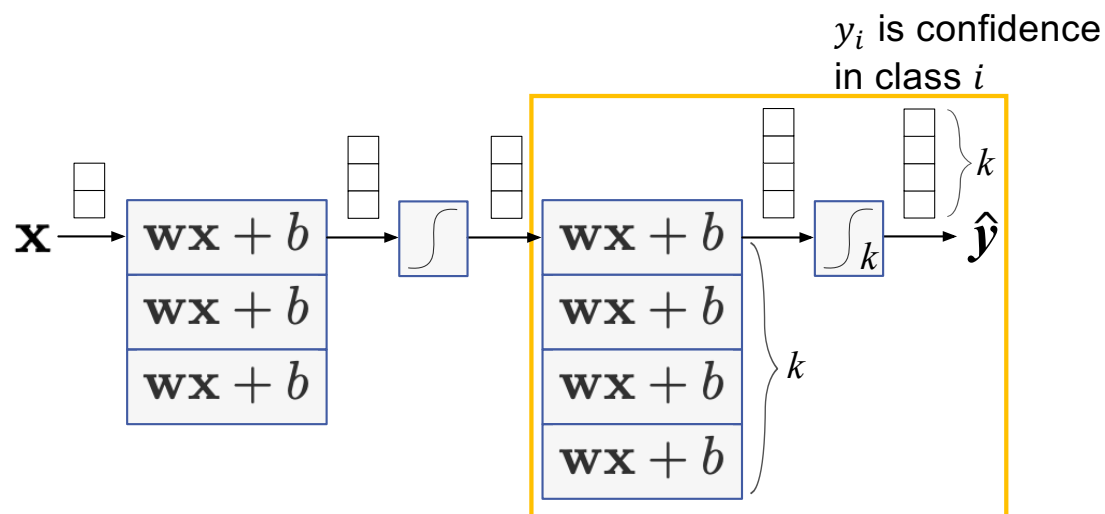


Multi-class classifiers

- **2-class problems:** final 1 neuron linear layer + sigmoid layer



- **k -class problems:** final k -neuron linear layer + softmax



k -class classifiers

- Instead of sigmoid, we use softmax function
- Instead of one neuron in last layer, we use k for k classes
- Last layer output: $\mathbf{z}^{[layer]} = W^{[layer]} \mathbf{x}^T + \mathbf{b}^{[layer]}$
- Vector of k probabilities as activation: $\mathbf{y} = \text{softmax}(\mathbf{z}^{[layer]})$
- Normalized probabilities of k class

$$\text{softmax}(\mathbf{z})_i = \frac{e^{z_i}}{\sum_{j=1}^k e^{z_j}}$$

Sample softmax computation

- For layer output vector \mathbf{z} :
$$\text{softmax}(\mathbf{z})_i = \frac{e^{z_i}}{\sum_{j=1}^k e^{z_j}}$$

```
z = np.array([0.1, 1, 5])  
np.exp(z)
```

```
array([ 1.10517092,  2.71828183, 148.4131591 ])
```

```
np.exp(z) / np.sum(np.exp(z))
```

```
array([0.00725956, 0.01785564, 0.9748848 ])
```

Training deep learning networks

What does training mean?

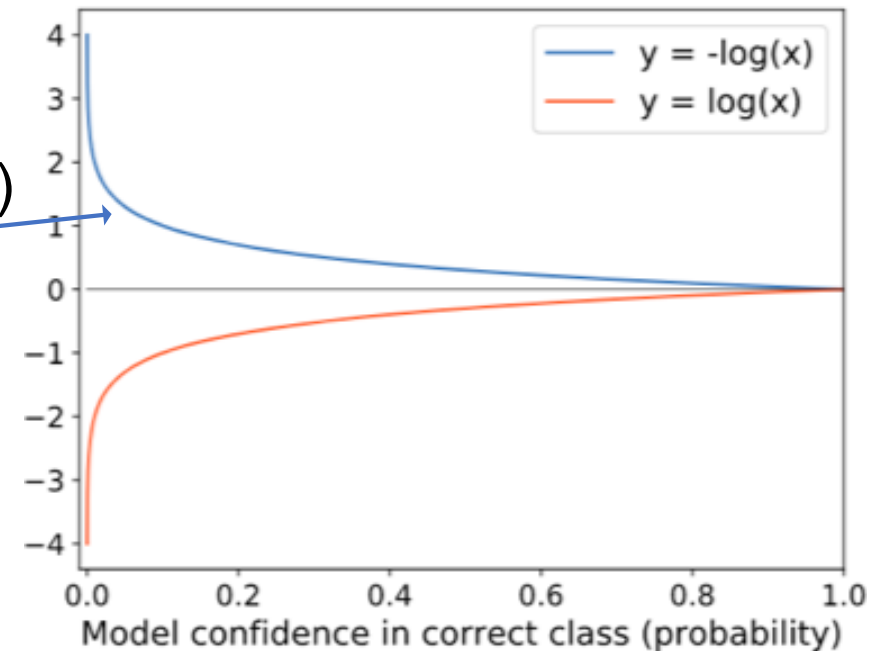
- Recall: testing an instance, a feature vector, means running that vector through the network and getting a prediction
- Prediction means computing a value using the model parameters; e.g., $\hat{y} = 3x + 2$ is a different model than $\hat{y} = .5x + 10$
- Find optimal (or good enough) model parameters as measured by a *loss* (cost) function
- Loss function measures the difference between model predictions and known targets
- We have huge search space (of parameters) and it is challenging to find parameters giving low loss

Loss functions

- **Regression:** typically mean squared error (MSE); should have smooth derivative, though mean absolute error works despite discontinuity (it's derivative is a V shape)
- **Classification:** log loss (also called cross entropy)
 - Penalizes very confident misclassifications strongly
 - Function of actual y and estimated probabilities, not predicted class
 - Perfect score is 0 log loss, imperfection gives unbounded scores
 - PyTorch log loss: `loss = cross_entropy(y_softmax, y_true)`
 - Predictions: `y_pred = argmax(y_softmax)`

Log loss

- $\text{loss} = \text{penalty}(p)$ if $y=1$ else $\text{penalty}(1-p)$
- Let $\text{penalty}(p) = -\log(p)$



- Two-class log loss:

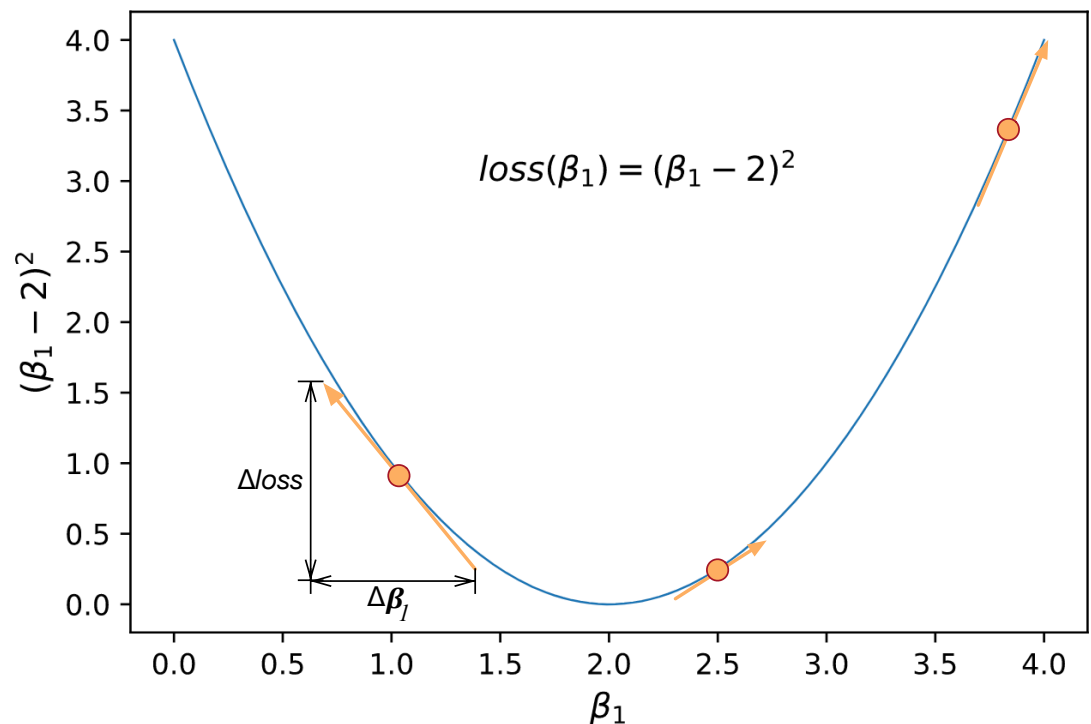
$$\text{loss} = -\frac{1}{n} \sum_{i=1}^n y_i \log(p) + (1 - y_i) \log(1 - p_i)$$

So log loss is average penalty where penalty is very high for confidence in wrong answer

Minimize loss with Gradient descent


- We use information about the loss function in the neighborhood of current parameters (here called β_i) to decide which direction shifts towards smaller loss
- Tweak parameters in that direction, amplified by a learning rate

```
while not_converged:  
     $\beta = \beta - \text{rate} * \text{gradient}(\beta)$ 
```



If learning rate is too high?

- We oscillate across valleys
- It can even diverge, exploding
- If too small, we don't make progress to min



Training process

1. Prepare data, normalize numeric variables
2. Split out at least a validation set from training set
3. Choose network architecture, appropriate loss function
4. Choose hyper-parameters, such as dropout rate
5. Choose a learning rate, number of epochs (passes through data)
6. Run training loop (until validation error goes up or num iterations)
7. Goto 3, 4, or 5 to tweak; iterate until good enough

Training loop

Regression

```
for epoch in range(nepochs):  
    y_train_pred = model(X_train)  
    loss = MSE(y_train, y_train_pred) # assume final sigmoid  
    update model parameters in direction of lower loss
```

Vectorized forward network pass



Classification

```
for epoch in range(nepochs):  
    y_train_pred = model(X_train)  
    loss = cross_entropy(y_train, y_train_pred) # assume softmax  
    update model parameters in direction of lower loss
```

What is vectorization?

- Use vectors not loops
- For torch/numpy arrays, we can use vector math instead of a loop:

$$\mathbf{c} = \mathbf{a} + \mathbf{b}$$

- Gives an opportunity to execute vector addition in parallel

```
for i in range(len(a)):  
    c[i] = a[i] + b[i]
```


	3	2	5	1	a
+	1	7	4	2	b
<hr/>					
	4	9	9	3	c



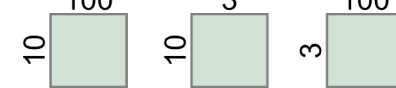
Vectorization in training loop

- Running one instance through network is how we think about it
- In practice, we send a subset or all X instances through the network in one go and compare all predictions to all y
- Instead of looping through instances, we pass all of X through to use matrix-matrix multiplies instead of matrix-vector multiplies

```
for epoch in range(nepochs):  
    for i in range(n):  
        y = model(X[i])  
    ...
```

$$\begin{array}{c} \frac{x}{3} \\ \text{yellow bar} \end{array} = \frac{X}{3} [i] \quad \begin{array}{c} \frac{y}{10} \\ \text{yellow bar} \end{array} = \frac{W}{3} @ \frac{x.T}{3}$$



```
for epoch in range(nepochs):  
    Y = model(X)  
    ...
```

$$\begin{array}{c} Y \\ 10 \times 100 \end{array} = \begin{array}{c} W \\ 10 \times 3 \end{array} @ \begin{array}{c} X.T \\ 3 \times 100 \end{array}$$


Assume $n=100$, $m=3$, $n_neurons=10$ in weight matrix W

Common train vs validation loss behavior

- DL networks have so many parameters, we can often get training error down to zero!
- But, we care about generalization
- Unfortunately, validation error often tracks away from training error as the number of epochs increases
- This model is clearly overfitting
- Need to use regularization to improve validation loss



Regularization techniques

- Get more training data; can try augmentation techniques (more data is likely to represent population distribution better)
- Reduce number of model parameters (i.e., simplify it) (not powerful enough to fit the noise)
- Add drop out layers (randomly kill some neurons)
- Weight decay (L2 regularization on model parameters)
- Early stopping, when validation error starts to go up (generally we choose model that yields the best validation error)
- Batch normalization has some small regularization effect (Force layer activation distributions to be 0-mean, variance 1)
- Stochastic gradient descent tends to land on better generalizations